

Lecture 3 Exercises

1. For a Bayesian linear regression model with prior on $p(w) = N(w|\mu, D^{-1})$ and q known, show that w_{Bayes} solves $(D + q\phi^T\phi)w_{Bayes} = q\phi^T t + D\mu$

To solve for bayes we take the known posterior probability:

$$-2\log p(w|t, X) = -2qt^T\phi w + qw^T\phi^T\phi w + (w - \mu)^T D(w - \mu) + const$$

To get the best estimate of w , we take the derivative of the posterior probability and set it equal to zero:

$$\partial - 2\log p(w|t, X) / \partial w = 0$$

To take the derivative of the first term:

$$\partial - 2qt^T\phi w / \partial w = -2qt^T\phi$$

The derivative of the second term uses the property $\partial x^T A x / \partial x = 2Ax$

$$\partial qw^T\phi^T\phi w / \partial w = 2q\phi^T\phi w$$

Expanding the third term $(w - \mu)^T D(w - \mu)$ gives:

$$= (w^T D - \mu^T D)(w - \mu)$$

$$= w^T D w - \mu^T D w - w^T D \mu + \mu^T D \mu$$

$$\partial (w^T D w - \mu^T D w - w^T D \mu + \mu^T D \mu) / \partial w = 2w^T D - 2\mu^T D$$

So we have:

$$0 = -2qt^T\phi + 2qw^T\phi^T\phi + 2w^T D - 2\mu^T D$$

Move negative terms to the LHS and divide both sides by 2:

$$qt^T\phi + \mu^T D = qw^T\phi^T\phi + w^T D$$

$$qw^T\phi^T\phi + w^T D = qt^T\phi + \mu^T D$$

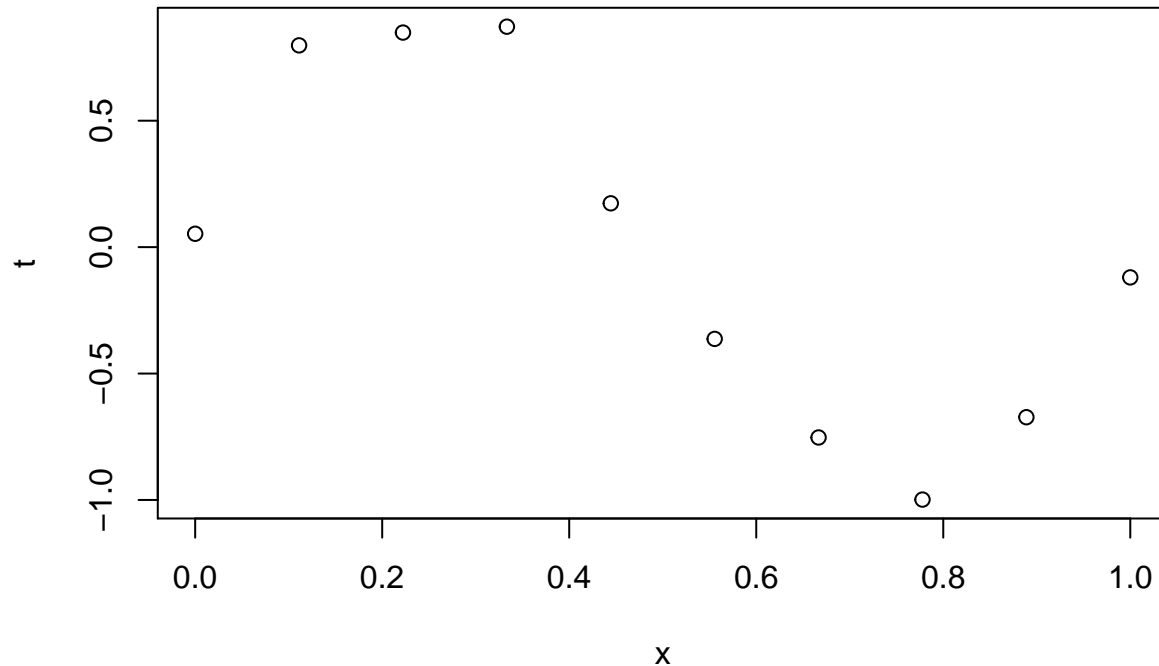
Take the transpose of both sides (note: $D^T = D$):

$$q\phi^T\phi w + Dw = q\phi^T t + D\mu$$

$$(D + q\phi^T\phi)w_{Bayes} = q\phi^T t + D\mu$$

2. Curve fitting (pt1) The aim is to learn a smooth function from the cloud of points stored in `curve_data.txt` using Bayesian linear regression models. In all that follows the prior $p(w) = N(w|0, \delta^{-1}I)$ is used and q, δ are constants specified by the user.

1. Plot the data



2.2 Write a function in R, called `phix` that takes as input a scalar x (the input in curve fitting), with values in $[0, 1]$, M the number of bases functions, and a categorical variable that specifies the type of basis used, and returns the vector of basis functions evaluated at x . Hence a call of the function `phix(0.3,4,"poly")` should return `c(1.0000, 0.3000, 0.0900, 0.0270, 0.0081)`. Code it up so that the option "poly" gives the polynomial bases and "Gauss" the Gaussian kernels with means μ_i equally spaced in $[0, 1]$, with $\mu = 0$ and $mu_M = 1$.

```
phix <- function(x, M, basis) {
  # Initialize vector to return
  u <- rep(0, M+1)
  # \mus for Gauss
  mus <- rep(0, M+1)
  interval <- 1 / M

  if (basis == 'poly') {
    # evaluate x at each value of p, e.g. x^0, x^1, ..., x^M
    for (p in 0:M) u[p+1] <- x^p
  } else if (basis == 'Gauss') {
    # Populate \mus with M intervals in [0,1]
    for (n in 1:M) mus[n+1] = mus[n] + interval
    # Generate Gaussian Kernels
    for (p in 0:M) u[p+1] <- exp(-(x - mus[p+1])^2/0.1)
  }

  u
}

# Test: round(phix(0.3,4,'poly'), digits = 4) == c(1, 0.3, 0.09, 0.027, 0.0081)
```

2.3. Write a function in R, called `post.params`, that takes as input the training data, M , the type of basis, the function `phix`, δ and q and returns the parameters of the posterior distribution, w_{Bayes} and Q .

```

library(MASS)

constructPhi <- function(features, M, basis_type) {
  phi <- matrix(nrow = length(features), ncol = M + 1)
  # Construct a phi matrix that is phi_x for each value of phi_x(x, M, 'Gauss') and append to phi matrix.
  for (i in 1:(length(features))) {
    phi[i,] = phi_x(features[i], M, basis_type)
  }
  phi
}

# Part 3: post.params
post.params <- function(training_data, M, basis_type, delta, q) {
  # regularization coefficient
  # training_data = curve_data
  # q = (1/0.1)^2
  # delta = 2.0
  # basis_type = 'Gauss'
  lambda = delta / q

  # Note: this requires the training data to have an x column
  xs = training_data$x
  # M = 9
  phi <- constructPhi(xs, M, basis_type)
  t = training_data$t

  # Slide 10 of Bayesian_regression.pdf
  w_Bayes = ginv(lambda * diag(M+1) + t(phi) %*% phi) %*% t(phi) %*% training_data$t

  D = delta * diag(M + 1)
  Q = D + q * (t(phi) %*% phi)

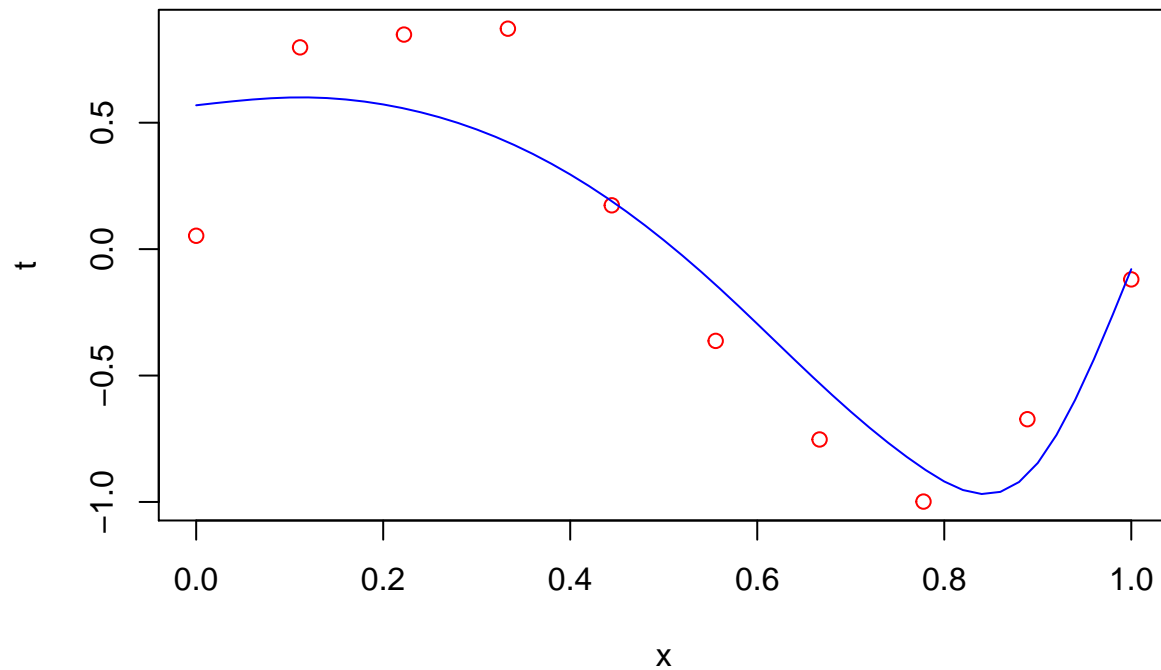
  # FIXME: RETURN Q?
  list(w_Bayes, Q)
}

M = 9
delta = 2.0
q = (1/0.1)^2

params <- post.params(curve_data, M, 'poly', delta, q)
w_Bayes <- params[[1]]
Q <- params[[2]]
phi <- constructPhi(curve_data$x, M, 'poly')
result <- phi %*% w_Bayes

# Part 4: Plot
plot(curve_data, col = 'red')
# Don't know how to smooth this out :(
library(splines)
lines(predict(splines::interpSpline(curve_data$x, result)), col = 'blue')

```



===== Roger's code:

```
library(ggplot2)
data <- read.table("curve_data.txt")

phix <- function(x, M, option) {
  phi <- rep(0,M)
  if (option == "poly") {
    for (i in 1:(M)) {
      phi[i] <- x**i
    }
  }
  if (option == "Gauss") {
    for (i in 0:(M-1)) {
      phi[i] <- exp(-((x-i/(M))**2)/0.1)
    }
  }
  phi
}

# Test: round(phix(0.3,4,'poly'), digits = 5) == c(0.3000, 0.0900, 0.0270, 0.0081)

post.params <- function(data, M, option, delta, q) {
  # Q = D + qPhiTPhi
  # wbyes = Q-1(qPhiTt + Dmu)

  phi = phix(data$x[1],M, option)
  for (i in 2:length(data$x)) {
    phi_ <- phix(data$x[i], M, option)
    phi = rbind(phi, phi_)
  }
  phi = cbind(rep(1,10),phi)
  phi
}
```

```

Q = diag(delta,M+1) + q*t(phi)%*%phi
#lambda = delta/q
w = solve(Q)%*%(q*t(phi)%*%data$t)
#w = solve(diag((2.0)/((1/0.1)**2),M+1)+t(phi)%*%phi)%*%t(phi)%*%data$t
t <- phi*w
return(list(Q,w,t))
}

answer <- post.params(data, 9, "poly", 2.0, (1/0.1)**2)

that <- answer[[3]]

p <- ggplot(cbind(data,that), aes(x=data$x, y=data$t)) + geom_point()
p + stat_smooth(method = "loess", formula = that ~ x, size = 1)

```

